DV1566- Project report

Group-49

Sai Nikhil Boyapati

Tabu Sravani Guduru

Problem Description:

Scalability is one of the major concerns while deploying the application. The server should be automatically able to scale in and out depending upon the no of requests are made. If it could not scale in when there are less no of requests, then it leads to the wastage of resources, and money and if a server could not scale out, it results in disappointed clients and bad server errors.

Implementation Technique:

The solution is implemented on AWS cloud. The services that we will be using to test the solution are EC2 instances, auto scaling group, and Load Balancer. The application we chose to deploy on the EC2 instance is a simple Django application.

Procedure:

The initial step is to login to the AWS management console and try to find the ec2 instance. We have setup the OS type to ubuntu, and we have launched the VMs. The default VPC and subnets are selected.

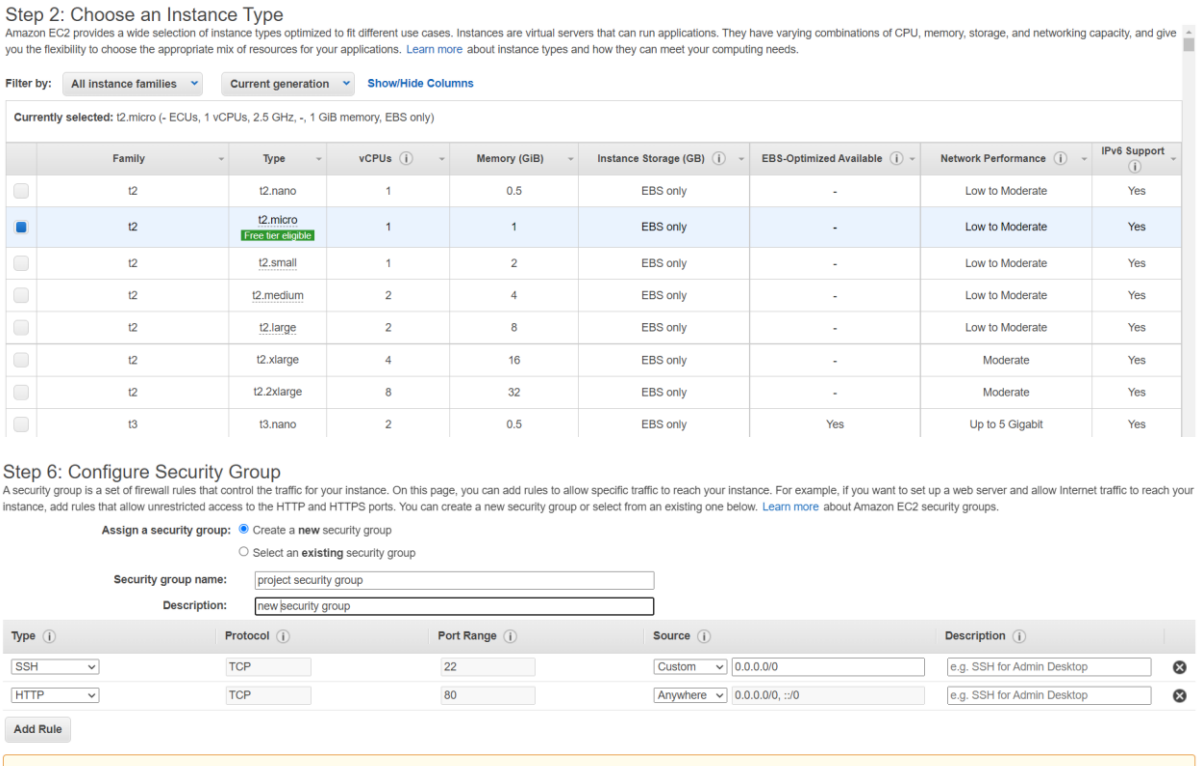

The security group is created according to our desired inbound rules.

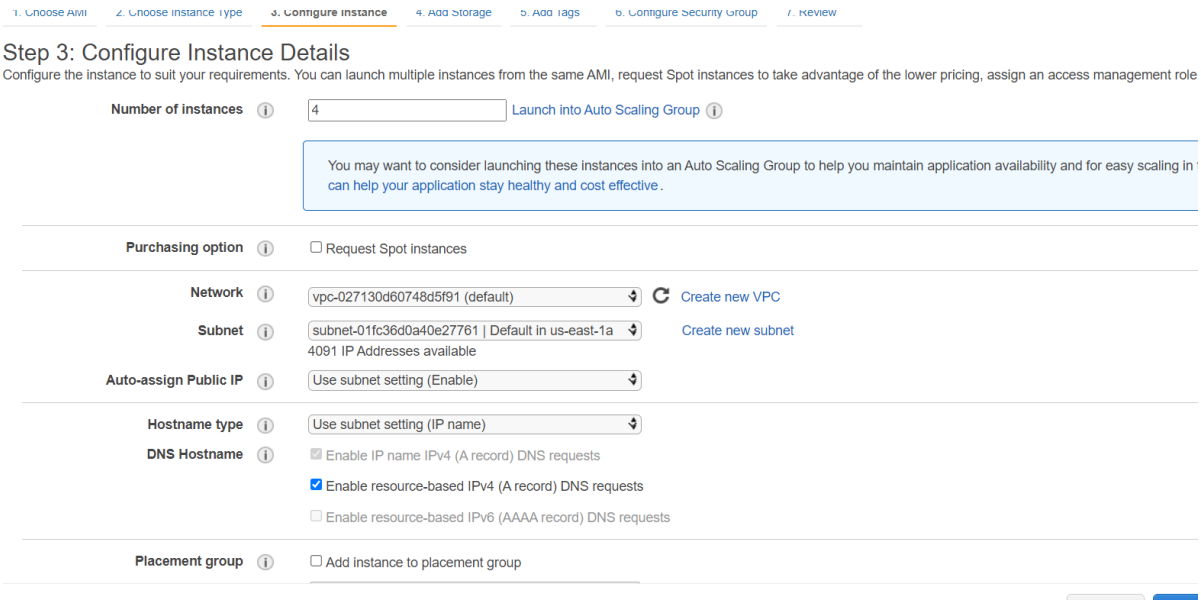

To enable the SSH connection, we have created the ssh key pair and it is added to the ~/.ssh folder and the permissions are set using chmod 600. The SSH connection is made to the linux.

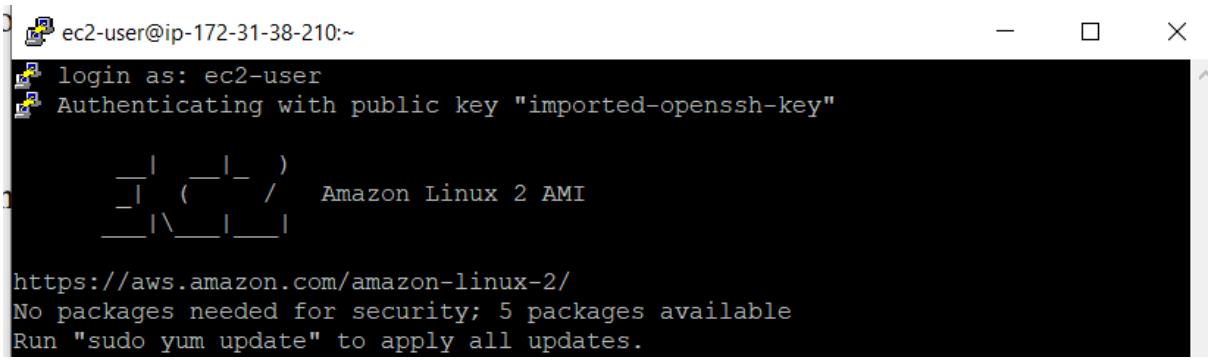

The next step is to set up the environment inside the instance to deploy the Django application. For that, we need to install some packages, the step-to-step procedure is given below. The name of the application that we set up is simple-django-app.

The

- Installing the required python and django packages in the ec2 instance.

The package for python is installed using the sudo apt update and the necessary environment is also set up. The commands that we used are

Sudo apt-get update && sudo apt-get upgrade

The Django is installed using pip3 install Django. The github application is cloned using git clone command. The nginx and gunicorn is installed using the commands

‘sudo apt-get install -y nginx’ && ‘pip3 install gunicorn’

The supervisor is installed. The supervisor allows the users to monitor and control a number of processes on UNIX-like operating systems.

‘sudo apt-get install supervisor’

Under the /etc/supervisor/conf.d. We need to create a file gunicorn.conf. The necessary lines has to be added to have a successful deployment. Later, test the configuration and allow nginx to read and restart the service nginx. The website is deployed and access the website using the public IP address of the EC2 instance.

Set up the autoscaling group to test the scalability feature:

Before setting up the auto scaling group, the image for an instance needs to be created.

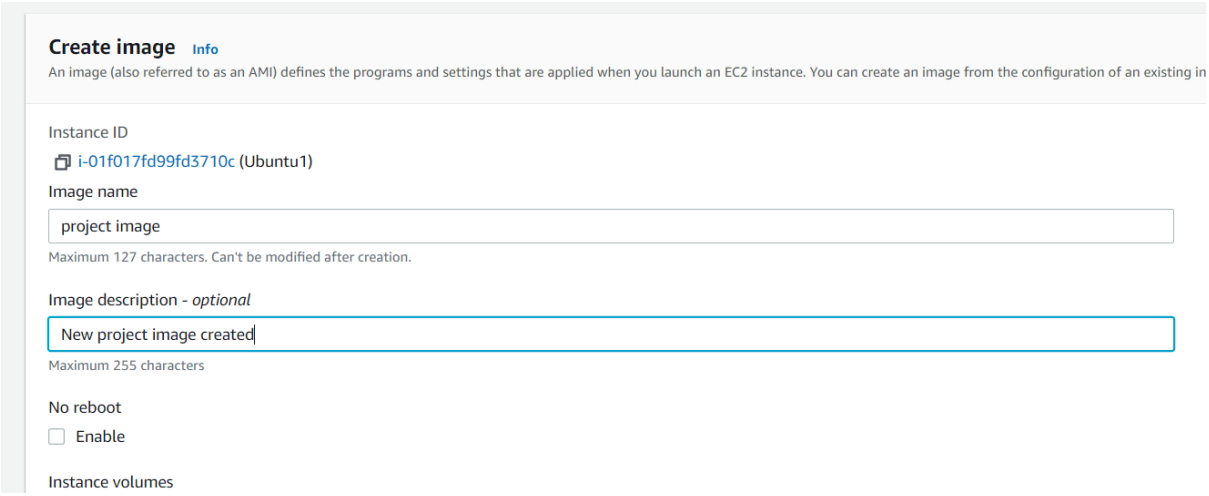

The load balancer and the target group is selected and the type of instance type and the desired number, the minimum and maximum capacity of the instances is selected. The keypair and the configuration is selected, and the process is shown below in the pictures.

- We have created a launch template named as launchproject and once creating is done and this template will be used in configuring the template which is used to creating auto scaling group.

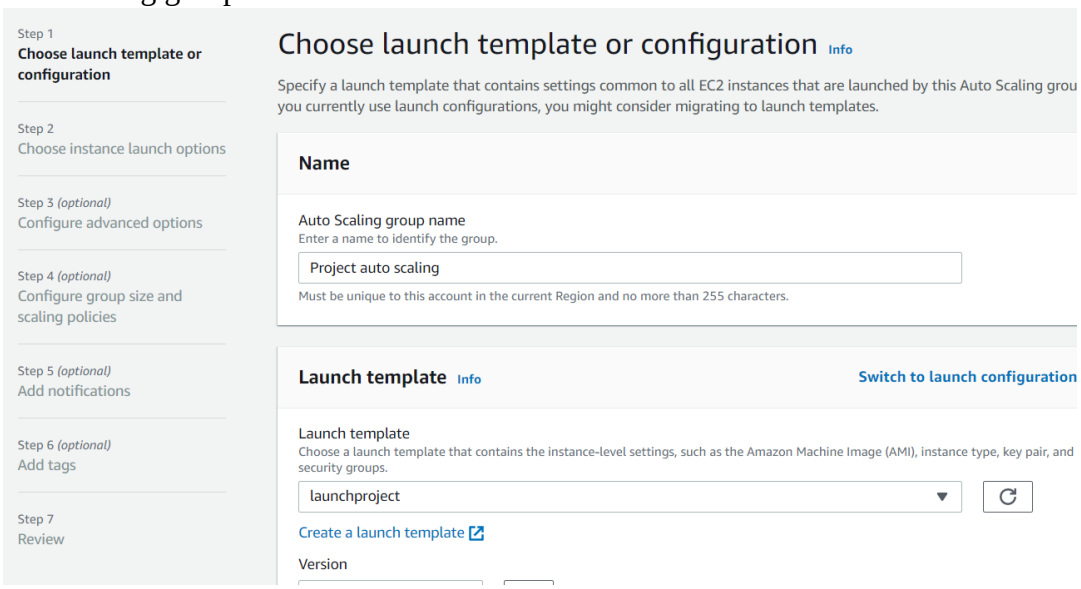

- In the figure below we can see that already created template is used while creating of auto scaling group process.

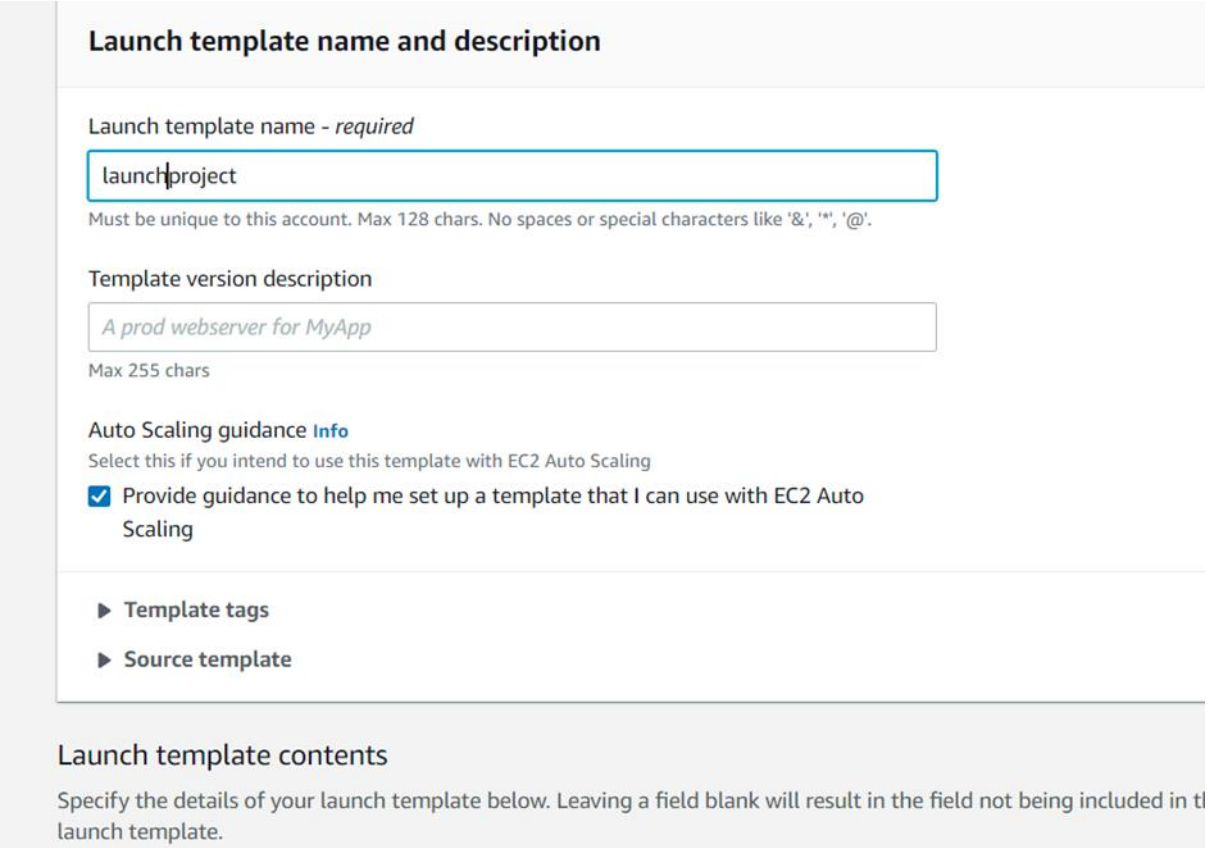

- In the launching process we have chosen instance type of t2.micro because it is free tier.

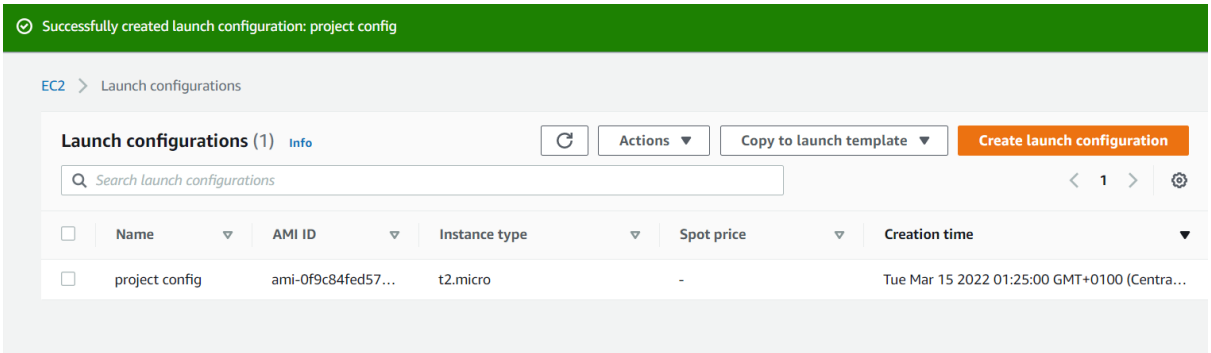

- In the next step we have chosen option existing load balancer which we already created target group which is named as projecttarget in the figure below.

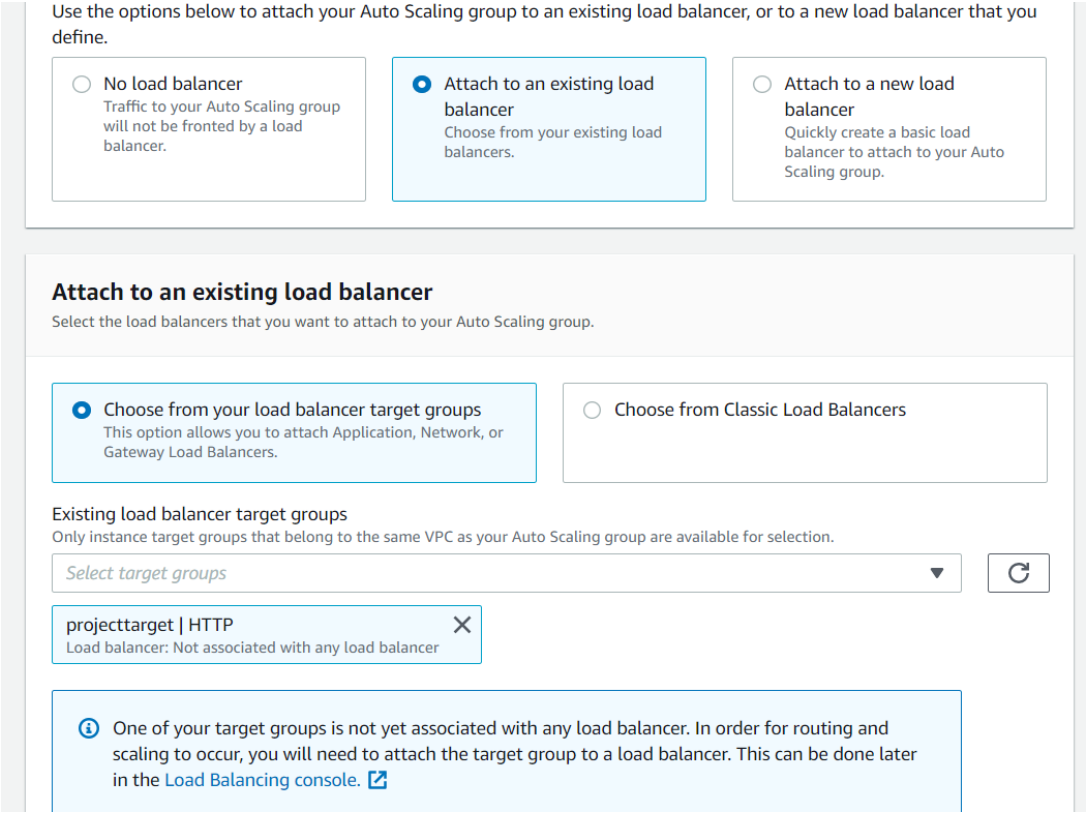

- After that we have chosen the instance which we want to register the target from four available instances.

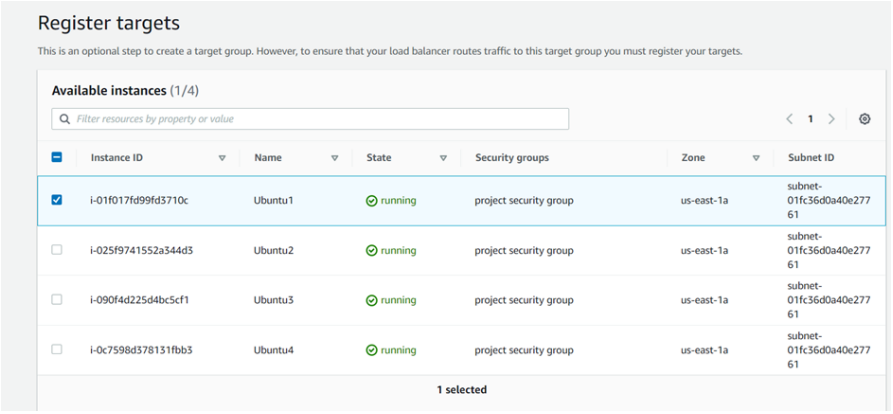

- In the last step we need capacities of instances for the application. In this project we have one instance for desired and minimum capacity and given three for maximum capacity as shown in figure below.

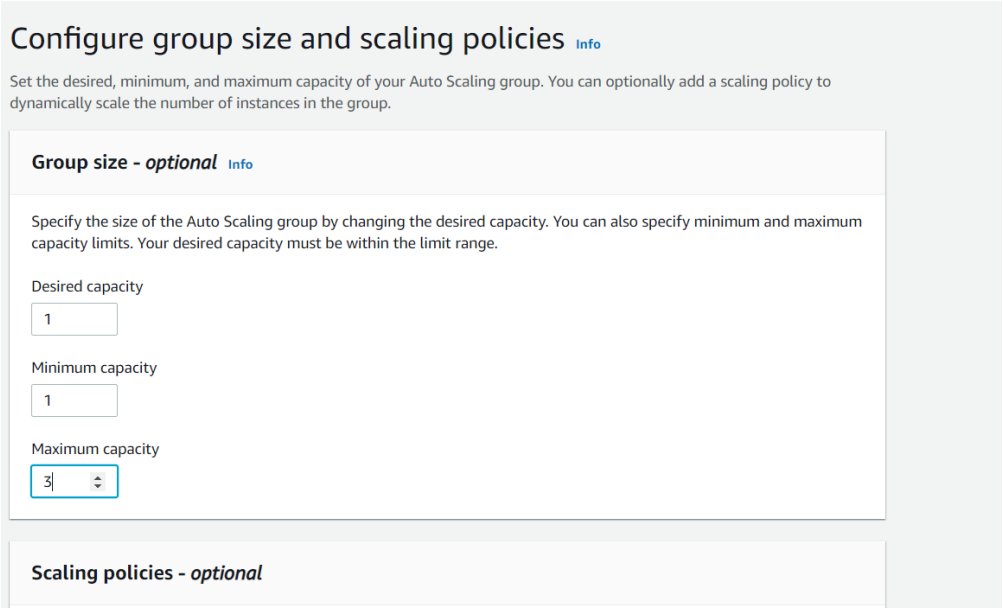

- The last step is to create the scaling the group. The following figure shows we have successfully created auto scaling group. In the scaling policy we have given CPU utilization upto 60%.

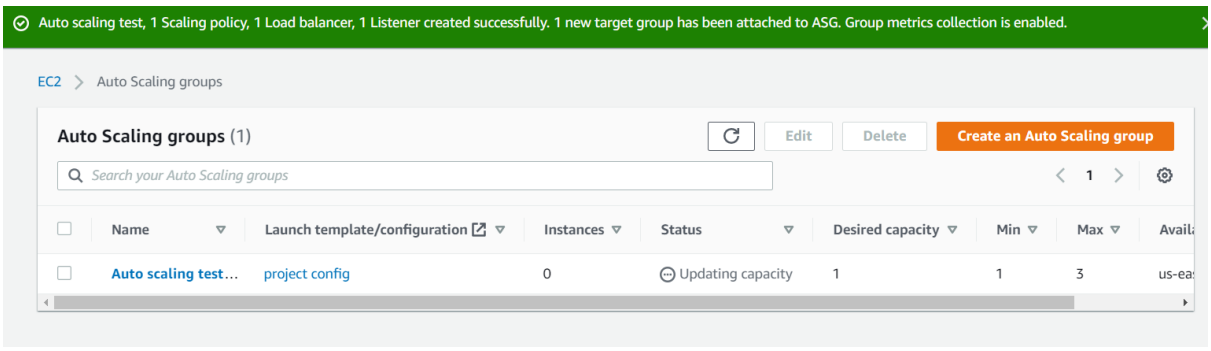

Testing through performing stress test:

After successfully, setting up the auto scaling group, stress test is performed to check if the auto scaling group is working or not. We have same type of test which we have used in lab2. To perform the stress test, the command that is given in the ubuntu is

ab -n 50000 -c ip-172-31-33-252.ec2.internal/index.html

```
ubuntu@ip-172-31-33-252:~$ ab -n 500000 -c 5 ip-172-31-33-252.ec2.internal/index.html
This is ApacheBench, Version 2.3 <$Revision: 1843412 $>
Copyright 1996 Adam Twiss, Zeus Technology Ltd, http://www.zeustech.net/
Licensed to The Apache Software Foundation, http://www.apache.org/

Benchmarking ip-172-31-33-252.ec2.internal (be patient)
Completed 50000 requests
Completed 100000 requests
Completed 150000 requests
Completed 200000 requests
Completed 250000 requests
Completed 300000 requests
Completed 350000 requests
Completed 400000 requests
Completed 450000 requests
Completed 500000 requests
Finished 500000 requests


Server Software:        nginx/1.18.0
Server Hostname:        ip-172-31-33-252.ec2.internal
Server Port:            80

Document Path:          /index.html
Document Length:        162 bytes

Concurrency Level:      5
Time taken for tests:   63.474 seconds
Complete requests:      500000
Failed requests:        0
Non-2xx responses:      500000
Total transferred:      160500000 bytes
HTML transferred:       81000000 bytes
Requests per second:    7877.24 [#/sec] (mean)
Time per request:       0.635 [ms] (mean)
Time per request:       0.127 [ms] (mean, across all concurrent requests)
Transfer rate:          2469.33 [Kbytes/sec] received

Connection Times (ms)
              min  mean[+/-sd] median   max
Connect:        0    0   0.1      0       7
Processing:     0    0   0.4      0      11
Waiting:        0    0   0.4      0      11
Total:          0    1   0.3      1      11

Percentage of the requests served within a certain time (ms)
  50%      1
  66%      1
```

After, the performance of the stress test, the CPU utilization has been raised a lot. It is shown in the picture given below. Increase in stress test resulted in increase in CPU utilization. Due to the auto scaling group, the instance has been created according to the autoscaling settings.

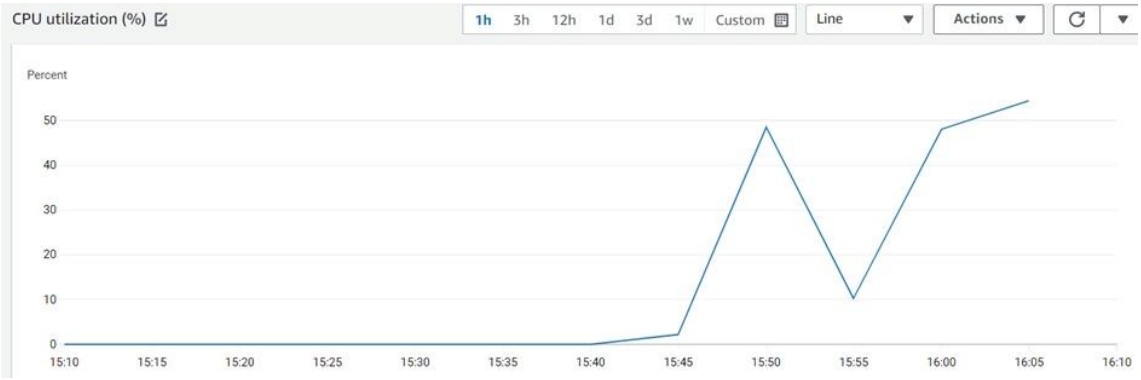

Conclusion:

The main aim of the problem description to test the scalability feature is verified. To perform this test, we have chosen a simple Django application. The environment is set up in the AWS cloud to deploy the application and to test the feature. We have incorporated the load balancer and the auto scaling group. After setting up the environment, we have performed stress test to check the CPU utilization, and after increase in the CPU utilization, the instance has been created. We have concluded that it is very feasible to set up the scalability of the instances through the environment provided by the AWS cloud provider.